

An Implementation of Round-Robin Scheduling for RIOT OS

Thanaphon Leonardi
University of Genova, DIBRIS
Via All'Opera Pia, 13
S4633433@studenti.unige.it

Dylan Lo Blundo
University of Genova, DIBRIS
Via All'Opera Pia, 13
S4623731@studenti.unige.it

Summary

A Round-Robin scheduling policy has been implemented with a time quantum of 0.5 seconds, replacing the previous priority-based approach.

Priorities have been maintained, however, in the thread structure, to maintain as much of the original operating system's structure as possible.

The **RRTester** application has been developed to verify the correct functionality of the new scheduler.

1 Introduction

A new chronologically ordered list has been created to keep track of all threads, including the idle thread (for which specific checks are made).

Due to the requirements of the **RRTester** application, the thread structure has been expanded to also include information on the service times of the various threads. To this end, the **xtimer** library has also been used, allowing the calculation of remaining service times and the switching of threads every quantum of time.

Finally, the *runqueue bitcache* and its relative system checks have been removed, to reduce superfluous work load.

2 Implementation Choices

2.1 File Changes

- *sched.c*
- *sched.h*
- *thread.c*
- *thread.h*

2.2 Tester Application

main.c

The **RRTester** application creates five threads, each with its own name and service time. To simulate the required workloads, the threads are put into busy waiting until the specified time has passed.

2.3 Thread Structure

thread.h

struct _thread

```
struct _thread {
    char *sp;                /**< thread's stack pointer    */
    thread_status_t status;  /**< thread's status        */
    uint8_t priority;        /**< thread's priority      */
    uint64_t service_time;   /**< thread's total executed time */
    uint64_t max_service_time; /**< thread's maximum lifetime */

    kernel_pid_t pid;        /**< thread's process id    */
}
```

service_time and *max_service_time* have been added to thread structure.

2.4 Thread Implementation

thread.c

thread_create()

The necessary changes have been made to keep the new thread list (*list_all_threads*) updated.

The idle thread, which behaves differently, is ignored by the Round-Robin queue until it is the last remaining one.

2.5 Scheduling Algorithm

sched.c

sched_run()

```
128     // Thread selection
129     if (active_thread) {
130         if (clist_count(&list_all_threads) == 1) {
131             next_thread = active_thread;
132         } else {
133             next_thread = container_of((&active_thread->rq_entry)->next,
134                                     thread_t, rq_entry);
135         }
136     } else {
137         // Only accept threads in PENDING state
138         next_thread = container_of((list_all_threads.next)->next, thread_t, rq_entry);
139     }
140
141     // Only accept threads in PENDING state
142     if (clist_count(&list_all_threads) > 1) {
143         while (next_thread->status != STATUS_PENDING || next_thread->priority == 15) {
144             next_thread = container_of((&next_thread->rq_entry)->next,
145                                     thread_t, rq_entry);
146
147             if (next_thread->status == STATUS_RUNNING)
148                 break;
149         }
150     }
151
152     if (active_thread) {
153         _unschedule(active_thread);
154     }
```

The thread selection algorithm scans through the queue until it finds the next *pending* thread (ready for execution). Appropriate checks have been implemented to ignore the idle thread, should there be other threads in the list.

2.6 Service Time Calculation

sched.c

sched_run()

```
167     if (service_time_start != 0) {
168         next_thread->service_time += xtimer_now64().ticks64 - service_time_start;
169     }
170
171     if (next_thread->priority != 7) { // if NOT main thread
172         service_time_start = xtimer_now64().ticks64;
173     }
174
175     sched_active_pid = next_thread->pid;
176     sched_active_thread = (volatile thread_t *)next_thread;
177
178     #ifdef DEVELHELP
179         if (!isIdle) {
180             printf("\nCurrently running thread: Thread %s\n", sched_active_thread->name);
181         }
182
183         if (next_thread->priority == 15) {
184             isIdle = true;
185         } else {
186             isIdle = false;
187         }
188     #endif
189
190     if(next_thread->max_service_time != 0) {
191         if (next_thread->max_service_time > next_thread->service_time) {
192             // Avoid overflow from subtracting unsigned numbers
193             float time = (float) ((next_thread->max_service_time - next_thread->service_time));
194             printf("Thread service time remaining: %.2fs\n", (float) time / 1000000);
195         } else {
196             printf("Thread service time remaining: 0.00s\n");
197         }
198     }
199
```

In this section, the calculation of the remaining service time; `service_time_start` represents the moment in which the last thread woke up, which is then subtracted from the current time (`timer_now64().ticks64`), returning the remaining time.

The first `if` ensures no calculation occurs when an initial time has not yet been assigned, while the second `if` disallows the main thread from being included in the calculation.

```
200     xtimer_set(&timer_rr, ROUND_ROBIN_TIME_QUANTUM);
```

This is the callback function: every time quantum, it executes the scheduling algorithm.